

Installing SW4 version 2.0

N. Anders Petersson*

Björn Sjögreen*

November 6, 2017

Contents

1	Introduction	1
2	Compilers and third party libraries	2
3	Unpacking the source code tar ball	4
4	Installing <i>SW4</i> with make	4
4.1	Basic compilation and linking of <i>SW4</i>	5
4.1.1	Mac machines	5
4.1.2	Linux machines	5
4.1.3	Using make	6
4.1.4	How do I setup the <code>make.inc</code> file?	7
4.2	Building <i>SW4</i> with proj.4 and/or efile support	7
4.3	Testing the <i>SW4</i> installation	8
5	Installing <i>SW4</i> with CMake	9
5.1	CMake Options	9
5.2	CTest	11
6	Installing the proj.4, euclid, and cencalvm packages	12
7	Disclaimer	13

1 Introduction

The sole purpose of this document is to describe the installation process of the seismic wave propagation code *SW4*. A comprehensive user's guide is provided in the report by Petersson and Sjögreen [1].

*Center for Applied Scientific Computing, Lawrence Livermore National Laboratory, PO Box 808, Livermore CA 94551. This work performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under contract DE-AC52-07NA27344. This is contribution LLNL-SM-741310.

2 Compilers and third party libraries

Before you can build *SW4* on your system, you must have

1. the **lapack** and **blas** libraries. These libraries provide basic linear algebra functionality and are pre-installed on many machines;
2. a MPI-2 library. This library provides support for message passing on parallel machines. Examples of open source implementations include Mpich-2 and OpenMPI. Note that the MPI-2 library must be installed even if you are only building *SW4* for a single core system.

To avoid incompatibility issues and linking problems, we recommend using the same compiler for the libraries as for *SW4*.

In order to use geographic projection and material models stored in the *rfile* format, you need to install the Proj4 before building *SW4*:

- The Proj.4 library, <http://trac.osgeo.org/proj>

If you also wish to use material models using the *efile* format, you also need to download and install two additional libraries:

- The Euclid e-tree library, <http://www-2.cs.cmu.edu/~euclid>
- The cencalvm library,
http://earthquake.usgs.gov/data/3dgeologic/cencalvm_doc/index.html

To simplify the build process, all libraries should be installed under the same directory, such that the library files (.so, .a, etc.) are in the lib sub-directory and the include files (.h) end up in the include sub-directory. See Section 6 for details.

Mac computers We recommend using the MacPorts package manager for installing the required compilers and libraries. Simply go to www.macports.org, and install macports on your system. With that in place, you can use the **port** command as follows

```
shell> sudo port install gcc72
shell> sudo port select --set gcc mp-gcc72
shell> sudo port install mpich-gcc7
shell> sudo port select --set mpi mpich-gcc7
```

Here, **gcc72** refers to version 7.2 of the Gnu compiler suite. Compiler versions are bound to change in the future, so the above commands will need to be modified accordingly. Before starting, make sure you install a version of gcc that is compatible with the MPI library package. The above example installs the **mpich** package using the gcc72 compilers, which includes a compatible Fortran compiler. Alternatively, you can use the **openmpi** package. Note that the **port select** commands are used to create shortcuts to the compilers and MPI environment. By using the above setup, the Gnu compilers can be accessed with **gcc** and **gfortran** commands, and the MPI compilers and execution environment are called **mpicxx**, **mpif90**, and **mpirun**, respectively.

The **lapack** and **blas** libraries are preinstalled on recent Macs and can be accessed using the **-framework Accelerate** link option. If that is not available or does not work on your machine, you can download **lapack** and **blas** from www.netlib.org.

Linux machines We here give detailed instructions for installing the third part libraries under 64 bit, Fedora Core 18 Linux. Other Linux variants use similar commands for installing software packages, but note that the package manager `yum` is specific to Fedora Core.

You need to have root privileges to install precompiled packages. Start by opening an xterm and set your user identity to root by the command

```
su -
```

Install the compilers by issuing the commands

```
yum install gcc
yum install gcc-c++
yum install gcc-gfortran
```

You install the `mpich2` library and include files with the command

```
yum install mpich2-devel
```

The executables and libraries are installed in `/usr/lib64/mpich2/bin` and `/usr/lib64/mpich2/lib` respectively. We suggest that you add `/usr/lib64/mpich2/bin` to your path. This is done with the command

```
export PATH=${PATH}:/usr/lib64/mpich2/bin
```

if your shell is `bash`. For `tcsh` users, the command is

```
setenv PATH ${PATH}:/usr/lib64/mpich2/bin
```

It is convenient to put the path setting command in your startup file, `.bashrc` or `.cshrc`, for `bash` or `csh/tcsh` respectively.

The `blas` and `lapack` libraries are installed with

```
yum install blas
yum install lapack
```

On our system, the libraries were installed in `/usr/lib64` as `libblas.so.3` and `liblapack.so.3`. For some unknown reason, the install program does not add links to these files with extension `.so`, which is necessary for the linker to find them. We must therefore add the links explicitly. If the libraries were installed elsewhere on your system, but you don't know where, you can find them with the following command:

```
find / -name "*blas*" -print
```

After locating the directory where the libraries reside (in this case `/usr/lib64`), we add links to the libraries with the commands:

```
cd /usr/lib64
ln -s libblas.so.3 libblas.so
ln -s liblapack.so.3 liblapack.so
```

Note that you need to have root privileges for this to work.

3 Unpacking the source code tar ball

To unpack the *SW4* source code, you place the file `sw4-v2.0.tgz` in the desired directory and issue the following command:

```
shell> tar xzf sw4-v2.0.tgz
```

As a result a new sub-directory named `sw4-v2.0` is created. It contains several files and sub-directories:

- `LICENSE.txt` License information.
- `INSTALL.txt` A link to this document.
- `README.txt` General information about *SW4*.
- `configs` Directory containing `make` configuration files.
- `src` C++ and Fortran source code of *SW4*.
- `tools` Matlab/Octave scripts for post processing and analysis.
- `pytest` Python script and input files for testing the *SW4* installation.
- `examples` Sample input files.
- `Makefile` Main makefile (don't change this file!).
- `CMakeLists.txt` CMake configuration file (don't change this file either!).
- `wave.txt` Text for printing the "SW4 Lives" banner at the end of a successful build.

4 Installing *SW4* with make

The classical way of building *SW4* uses `make`. We recommend using GNU `make`, sometimes called `gmake`. You can check the version of `make` on your system with the command

```
shell> make -v
```

If you don't have GNU `make` installed on your system, you can obtain it from www.gnu.org.

We have built *SW4* and its supporting libraries on Intel based laptops and desktops running LINUX and OSX. It has also been built on several supercomputers such as the Intel machines `cab`, `quartz` (at LLNL) and `edison`, `cori` (at LBNL), as well as the IBM BGQ machine `vulcan` at LLNL. We have successfully used the following versions of Gnu, Intel, and IBM compilers:

Gnu:	<code>g++/gcc/gfortran</code>	versions 4.5 to 7.2
Intel:	<code>icpc/icc/ifort</code>	versions 16.0 to 18.0
IBM Blue Gene:	<code>xlcc/xlc/xlf</code>	versions 12.1 to 14.1

SW4 uses the message passing interface (MPI) standard (MPI-2 to be specific) for communication on parallel distributed memory machines. Note that the MPI library often includes wrappers for compiling, linking, and running of MPI programs. For example, the `mpich2` package build wrappers for the underlying C++ and Fortran compilers called `mpicxx` and `mpif90`, as well as the `mpirun` script. We highly recommend using these programs for compiling, linking, and running *SW4*.

4.1 Basic compilation and linking of *SW4*

The best way of getting started is to first build *SW4* without the proj.4 and cencalvm libraries. This process should be very straight forward and the resulting *SW4* executable will support all commands except `rfile`, `efile` and the `proj/ellps` options in the `grid` command. If you need to use these options, you can always recompile *SW4* after the proj.4 and cencalvm libraries have been installed. See § 6 for details.

The basic build process is controlled by the environmental variables `FC`, `CXX`, `EXTRA_FORT_FLAGS`, `EXTRA_CXX_FLAGS`, and `EXTRA_LINK_FLAGS`. These variables should hold the names of the Fortran and C++ compilers, and any extra options that should be passed to the compilers and linker. The easiest way of assigning these variables is by creating a file in the `configs` directory called `make.inc`. The `Makefile` will look for this file and read it if it is available. There are several examples in the `configs` directory, e.g. `make.osx` for Macs and `make.linux` for Linux machines. You should copy one of these files to your own `make.inc` and edit it as needed.

4.1.1 Mac machines

If you are on a Mac, you could copy the setup from `make.osx`,

```
shell> cd configs
shell> cp make.osx make.inc
shell> cat make.inc
etree = no
proj = no
FC = mpif90
CXX = mpicxx
EXTRA_FORT_FLAGS =
EXTRA_LINK_FLAGS = -framework Accelerate -L/opt/local/lib/gcc72 -lgfortran
```

In this case, the `blas` and `lapack` libraries are assumed to be provided by the `-framework Accelerate` option. The `libgfortran` library is located in the directory `/opt/local/lib/gcc72`, which is where `macports` currently installs it.

4.1.2 Linux machines

If you are on a Linux machine, we suggest you copy the configuration options from `make.linux`,

```
shell> cd configs
shell> cp make.linux make.inc
shell> cat make.inc
FC = gfortran
CXX = mpicxx
EXTRA_LINK_FLAGS = -L/usr/lib64 -llapack -lblas -lgfortran
```

This setup assumes that the `blas` and `lapack` libraries are located in the `/usr/lib64` directory. In the case of Fedora Core 18, we needed to set the link flag variable to

```
EXTRA_LINK_FLAGS = -Wl,-rpath=/usr/lib64/mpich2/lib -llapack -lblas -lgfortran
```

4.1.3 Using make

You build *SW4* with the "make" command from the main directory.

```
shell> cd /enter/your/path/sw4-v2.0
shell> make
```

If all goes well, you will see the SW4 Lives banner on your screen after the compilation and linking has completed,

```
''_.,_.,.-''_.,_.,.=''_.,_.,.-''_.,_.,.=''_.,_.,.-''_.,_.,.=''
```

```

  -----
 /  _  \  /  _  \  /  _  \  /  _  \  /  _  \  /  _  \
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
\  _  /  \  _  /  \  _  /  \  _  /  \  _  /  \  _  /
/  _  \  /  _  \  /  _  \  /  _  \  /  _  \  /  _  \
\  _  /  \  _  /  \  _  /  \  _  /  \  _  /  \  _  /
  -----

```

```

  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
|  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _
|  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _
|  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _
|  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _  _

```

```
''_.,_.,.-''_.,_.,.=''_.,_.,.-''_.,_.,.=''_.,_.,.-''_.,_.,.=''
```

By default, **make** builds an optimized **sw4** executable. It is located in

```
/enter/your/path/sw4-v2.0/optimize/sw4
```

You can also build an executable with debugging symbols by adding the **debug=yes** option to **make**,

```
shell> cd /enter/your/path/sw4-v2.0
shell> make debug=yes
```

In this case, the executable will be located in

```
/enter/your/path/sw4-v2.0/debug/sw4
```

It can be convenient to add the corresponding directory to your **PATH** environment variable. This can be accomplished by modifying your shell configuration file, e.g. `~/.cshrc` if you are using C-shell.

4.1.4 How do I setup the `make.inc` file?

The input file for `make` is

```
sw4-v2.0/Makefile
```

Do *not* change this `Makefile`. It should only be necessary to edit your configuration file, that is, `/my/path/sw4-v2.0/configs/make.inc`

Note that you must create this file, for example by copying one of the `make.xyz` files in the same directory. The `make.inc` file holds all information that is particular for your system, such as the name of the compilers, the location of the third party libraries, and any extra arguments that should be passed to the compiler or linker. This file also tells `make` whether or not the `cencalvm` and `proj.4` libraries are available and where they are located.

The following `make.inc` file includes all configurable options:

```
etree = no
proj = no
SW4ROOT = /Users/petersson1

CXX = mpicxx
FC  = mpif77

EXTRA_CXX_FLAGS = -DUSING_MPI
EXTRA_FORT_FLAGS = -fno-underscoring
EXTRA_LINK_FLAGS = -framework vecLib
```

The `etree` variable should be set to `yes` or `no`, to indicate whether or not the `cencalvm` and related libraries are available. The `SW4ROOT` variable is only used when `etree=yes`. The `CXX` and `FC` variables should be set to the names of the C++ and Fortran compilers, respectively. Finally, the `EXTRA_CXX_FLAGS`, `EXTRA_FORT_FLAGS`, and `EXTRA_LINK_FLAGS` variables should contain any additional arguments that need to be passed to the C++ compiler, Fortran compiler, or linker, on your system.

4.2 Building *SW4* with `proj.4` and/or `efile` support

The installation of the `proj.4`, `euclid`, and `cencalvm` libraries is discussed in Section 6. Note that the `proj.4` library enables the more advanced geographical mapping keywords in the `grid` command and is also required by the `rfile` command. To enable the `efile` command, you have to also install the `euclid` and `cencalvm` libraries. Note that the latter two libraries are only needed by the `efile` command. If you are not planning on using that command, there is no need to install those libraries. This is a change from *SW4* version 1.0.

Once you have successfully installed the `proj.4`, and optionally the `euclid` and `cencalvm` libraries, it should be easy to re-configure *SW4* to use them. Simply edit your configuration file (`make.inc`) by adding two lines to the top of the file, setting the `etree` keyword to `yes` or `no`, as appropriate.

```
proj = yes
etree = no
SW4ROOT = /thirdparty/basedir
...
```

You then need to re-compile *SW4*. Go to the *SW4* main directory, clean out the previous object files and executable, and re-run make:

```
shell> cd /my/installation/dir/sw4-v2.0
shell> make clean
shell> make
```

If all goes well, the “SW4 lives” banner is shown after the make command is completed. As before, the *sw4* executable will be located in the *optimize* or *debug* directories.

4.3 Testing the *SW4* installation

The *SW4* source code distribution includes a python(3) script for running several tests and checking the solutions against previously verified results. Note that the same set of tests can be performed when *SW4* is built with CMake, see Section 5.2.

After *SW4* has been built with *make*, go to the *pytest* directory and run *test_sw4.py*. If the *sw4* executable resides in the *optimize* directory, you can run the basic tests by doing:

```
shell> cd pytest
shell> ./test_sw4.py
```

If all goes well, you should see the following output:

```
>shell test_sw4.py
Running all tests for level 0 ...
Starting test # 1 in directory: meshrefine with input file: refine-el-1.in
Test # 1 Input file: refine-el-1.in PASSED
Starting test # 2 in directory: meshrefine with input file: refine-att-1.in
Test # 2 Input file: refine-att-1.in PASSED
...
Starting test # 12 in directory: lamb with input file: lamb-1.in
Test # 12 Input file: lamb-1.in PASSED
Out of 12 tests, 12 passed and 0 failed.
```

Some aspects of the testing can be modified by providing command line arguments to *test_sw4.py*. For a complete list of options do *test_sw4.py --help*, which currently give the output:

```
shell> ./test_sw4.py --help
usage: test_sw4.py [-h] [-v] [-l {0,1,2}] [-m MPITASKS] [-d SW4_EXE_DIR]
```

optional arguments:

```
-h, --help            show this help message and exit
-v, --verbose          increase output verbosity
-l {0,1,2}, --level {0,1,2}
                        testing level
-m MPITASKS, --mpitasks MPITASKS
                        number of mpi tasks
-d SW4_EXE_DIR, --sw4_exe_dir SW4_EXE_DIR
                        name of directory for sw4 executable
```

Note that the directory name for the *sw4* executable should be given relative to the main *sw4* directory.

5 Installing *SW4* with CMake

SW4 can also be built with CMake. Compared to using regular `make`, this build process is easier to use because it is fully automated. However, it gives the user less control of which compilers, linker, and libraries to use. Similar to using regular `make`, the *SW4* CMake configuration allows automated correctness testing of the installation. The test runs the same set of cases as the `test_sw4.py` script in the `pytest` directory, see Section 5.2 for details.

To use CMake, navigate to the top `sw4` directory and run the following commands:

```
shell> mkdir build
shell> cd build
shell> cmake [options] ..
shell> make
```

The two dots after `cmake [options]` are essential and instructs it to look in the parent directory for the `CMakeLists.txt` file.

The `cmake` command searches for the necessary libraries and other dependencies, then creates makefiles that are appropriate for your system. You then run `make` to compile and link *SW4* using these makefiles. For details about the exact commands being used in compilation, run `make VERBOSE=1`. Once *SW4* has been successfully built, you will see the “SW4 Lives!” banner on the screen.

NOTE: `cmake` puts the `sw4` executable in the `build/bin` directory. This is different from regular `make`, which puts the executable in the `optimize` directory.

NOTE: If you want to rebuild `sw4` with a new set of options, you can force `cmake` to start from scratch by removing the file `CMakeCache.txt` in the `build` directory. Another way is to remove all files in the `build` directory.

5.1 CMake Options

CMake provides several options to allow customized configuration of *SW4*. To use any option, add `-D<option>=<value>` to the options in the `cmake` command. For example:

```
cmake -DTESTING_LEVEL=1 -DCMAKE_BUILD_TYPE=Debug ..
```

configures *SW4* with testing level 1, to be compiled with debugging symbols in the object files. A list of options is shown in the table below.

Option	Default	Details
PROJ4_HOME	(none)	The path to the Proj.4 installation to use when compiling <i>SW4</i> .
CENCALVM_HOME	(none)	The path to the cencalvm installation to use when compiling <i>SW4</i> .
CMAKE_BUILD_TYPE	Release	The type of build to setup. Can be either <code>Debug</code> , <code>Release</code> , or <code>RelWithDebInfo</code> . This affects the type of optimization and debug flags used in compiling <i>SW4</i> .
TESTING_LEVEL	0	Specifies the testing level for automated tests. Level 0 corresponds to tests that run in roughly a minute or less (7 total), level 1 to tests that run in roughly 10 minutes or less (13 total) and level 2 to tests that may require up to an hour or more (17 total).
MPI_NUM_TEST_PROCS	4	Number of MPI processes to use in tests. Generally using more processes will result in the tests finishing faster, but there is no point exceeding the number of available cores on your system. We strongly recommend at least 8 processes if <code>TESTING_LEVEL</code> is 1 or higher.
MPIEXEC	mpirun	UNIX command for running an MPI application.
MPIEXEC_NUMPROC_FLAG	-np	MPI command option for specifying the number of processes.
MPIEXEC_PREFLAGS	(none)	Extra MPI command option.

Modifying the MPI execution commands. By default, `mpirun` is used to start parallel runs when you do `make test`. However, on Livermore computing (LC) machines the command for running MPI programs is `srun`, not `mpirun`. Also, the flag for specifying the number of processors is different, and you must give an additional flag for running interactive jobs on the debug partition. For example, you would say

```
srun -ppdebug -n 128 sw4 inputfile.in
```

to run on the debug partition using 128 cores. To modify the default MPI execution program and other runtime parameters, the variables `MPIEXEC`, `MPIEXEC_NUMPROC_FLAG`, and `MPIEXEC_PREFLAGS` can be set as in the following example:

```
cmake -DTESTING_LEVEL=2 -DMPI_NUM_TEST_PROCS=128 -DMPIEXEC=/usr/bin/srun \
      -DMPIEXEC_NUMPROC_FLAG=-n -DMPIEXEC_PREFLAGS=-ppdebug ..
```

After the `proj.4`, `euclid` and `cencalvm` libraries have been installed (see next section), you need to tell `cmake` where to find them. On the LC-machines, all three libraries are currently installed under `/usr/apps/wpp`, and you can use the following command options to configure `sw4`:

```
cmake -DTESTING_LEVEL=2 -DMPI_NUM_TEST_PROCS=36 -DMPIEXEC=/usr/bin/srun \
      -DMPIEXEC_NUMPROC_FLAG=-n -DMPIEXEC_PREFLAGS=-ppdebug \
      -DPROJ4_HOME=/usr/apps/wpp -DCENCALVM_HOME=/usr/apps/wpp ..
```

To verify that `cmake` actually found the libraries, pay attention to the following lines of the output from the `cmake` command:

```
...
-- Found PROJ4: /usr/apps/wpp/lib/libproj.so
-- Found CENCALVM: /usr/apps/wpp/lib/libcencalvm.so;/usr/apps/wpp/lib/libetree.so
...
```

Sometimes CMake doesn't pick up the correct compiler. Say, for example that the C++ compiler on your system is called `mpicxx` and the Fortran compiler is `mpiifort`. You can tell `cmake` to use those compilers by setting the following environment variables *before* running `cmake` (assuming a `csh` shell),

```
> setenv CXX mpicxx
> setenv FC mpiifort
```

5.2 CTest

The *SW4* CMake configuration includes several test cases that are used to verify the correctness of the *SW4* installation. Each test consists of two parts. First it runs a case using an input file in the `pytest` directory. Secondly, it checks that the results are within a reasonable error tolerance from previously recorded results.

To run the tests, use either the command `make test` or `ctest` as follows:

```
build > ctest
Test project /Users/petersson1/src/sw4-cig/build
  Start 1: Run_twilight/flat-twi-1
1/24 Test #1: Run_twilight/flat-twi-1 ..... Passed    0.49 sec
  Start 2: Check_Result_twilight/flat-twi-1
2/24 Test #2: Check_Result_twilight/flat-twi-1 ..... Passed    0.03 sec
  Start 3: Run_twilight/flat-twi-2
...
  Start 23: Run_pointsource/pointsource-sg-1
23/24 Test #23: Run_pointsource/pointsource-sg-1 ..... Passed   89.56 sec
  Start 24: Check_Result_pointsource/pointsource-sg-1
24/24 Test #24: Check_Result_pointsource/pointsource-sg-1 ... Passed    0.03 sec

100\% tests passed, 0 tests failed out of 24

Total Test time (real) = 230.91 sec
```

You can run tests selectively using `ctest -R <regex>`, for example:

```
build > ctest -R meshrefine
Test project /Users/petersson1/src/sw4-cig/build
  Start 15: Run_meshrefine/refine-el-1
1/6 Test #15: Run_meshrefine/refine-el-1 ..... Passed   25.61 sec
  Start 16: Check_Result_meshrefine/refine-el-1
2/6 Test #16: Check_Result_meshrefine/refine-el-1 ..... Passed    0.03 sec
```

```

    Start 17: Run_meshrefine/refine-att-1
3/6 Test #17: Run_meshrefine/refine-att-1 ..... Passed    22.00 sec
    Start 18: Check_Result_meshrefine/refine-att-1
4/6 Test #18: Check_Result_meshrefine/refine-att-1 ..... Passed     0.03 sec
    Start 19: Run_meshrefine/refine-att-2nd-1
5/6 Test #19: Run_meshrefine/refine-att-2nd-1 ..... Passed    17.63 sec
    Start 20: Check_Result_meshrefine/refine-att-2nd-1
6/6 Test #20: Check_Result_meshrefine/refine-att-2nd-1 ... Passed     0.03 sec

```

100% tests passed, 0 tests failed out of 6

Total Test time (real) = 65.35 sec

If a test fails you can check the details in the output log at `Testing/Temporary/LastTest.log`.

6 Installing the proj.4, euclid, and cencalvm packages

If you are only interested in using the advanced geographical mapping options of the `grid` command, or the `rfile` command, you only need to install the proj.4 package.

The following instructions describe how to install all three packages. For simplicity all packages are installed under the same top directory. If you are using `cmake`, you may optionally put the proj.4 package in a separate directory. In the following we shall assume that all packages are installed under the same top directory, and that you assign the name of that directory to the environment variable `SW4ROOT`. When you are finished installing the packages, the corresponding include and library files should be in the sub-directories `${SW4ROOT}/include` and `${SW4ROOT}/lib`, respectively.

The cencalvm library was developed by Brad Aagaard at USGS. Instructions for building the cencalvm library as well as downloading the Etree database files for Northern California, can currently be downloaded from

http://earthquake.usgs.gov/data/3dgeologic/cencalvm_doc/INSTALL.html

The installation process for cencalvm, which is outlined below, is described in detail on the above web page. Note that cencalvm relies on both the euclid and the proj.4 libraries.

The euclid library must be installed manually by explicitly copying all include files to the include directory and all libraries to the lib directory,

```

shell> cd euclid3-1.2/libsrc
shell> make
shell> cp *.h ${SW4ROOT}/include
shell> cp libetree.* ${SW4ROOT}/lib

```

The proj4 library should be configured to be installed in `${SW4ROOT}`. This is accomplished by

```

shell> cd proj-4.7.0
shell> configure --prefix=${SW4ROOT}
shell> make
shell> make install

```

We remark that the `proj.4` library can alternatively be installed using `macports` (if you are on a Mac OSX machine).

The `cencalvm` library should also be configured to be installed in `${SW4ROOT}`. You also have to help the configure script finding the include and library files for the `proj.4` and `etree` libraries,

```
shell> cd cencalvm-0.6.5
shell> configure --prefix=${SW4ROOT} CPPFLAGS="-I${SW4ROOT}/include" \
               LDFLAGS="-L${SW4ROOT}/lib"
shell> make
shell> make install
```

To verify that the libraries have been installed properly, you should go to the `SW4ROOT` directory and list the `lib` sub-directory (`cd ${SW4ROOT}; ls lib`). You should see the following files (on Mac OSX machines, the `.so` extension is replaced by `.dylib`):

```
shell> cd ${SW4ROOT}
shell> ls lib
libetree.so libetree.a
libproj.so libproj.a libproj.la
libcencalvm.a libcencalvm.la libcencalvm.so
```

Furthermore, if you list the include sub-directory, you should see include files such as

```
shell> cd ${SW4ROOT} %$
shell> ls include
btree.h etree.h etree_inttypes.h
nad_list.h projects.h proj_api.h
cencalvm
```

Note that the include files for `cencalvm` are in the sub-directory with the same name.

7 Disclaimer

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.

References

- [1] N. A. Petersson and B. Sjögren. User's guide to SW4, version 2.0. Technical Report LLNL-SM-abcdef, Lawrence Livermore National Laboratory, 2017. (Source code available from geodynamics.org/cig).